




# Machine Learning

2025 (ML-2025)

Lecture 1. Intro to ML

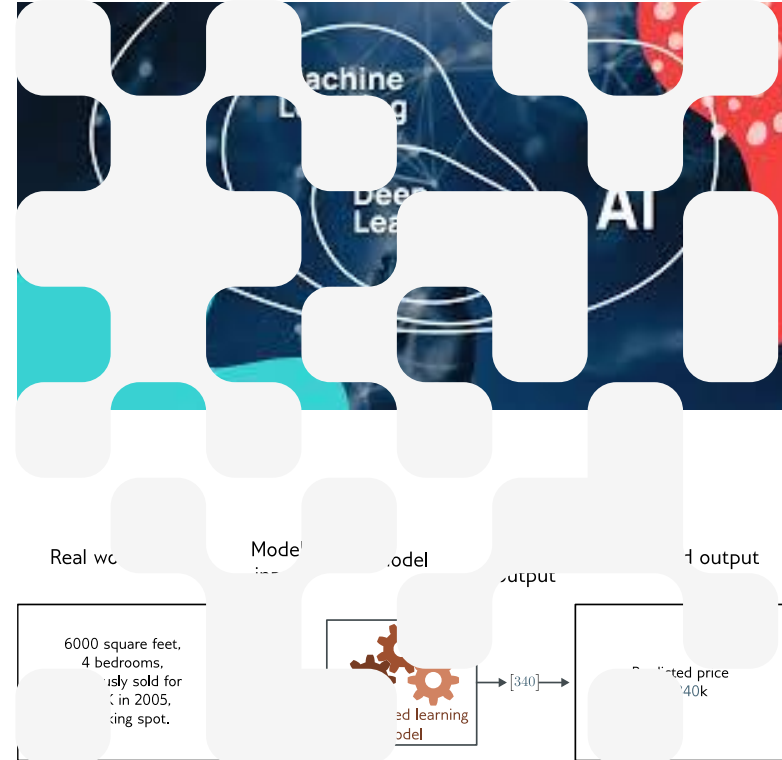
by Alexei Kornaev, Dr. Sc., Assoc. Prof., Robotics and CV, [Innopolis University](#)  
Researcher at the RC for AI, [National RC for Oncology n.a. NN Blokhin](#)



# Agenda

- I. **Logistics**
- II. ML Overview:  $ML = E + T + P$
- III. Experience (E) in ML
- IV. ML Tasks (T)
- V. Performance measure (P) in ML

All models are wrong, but some are useful.  
/George Box/



## Course objectives

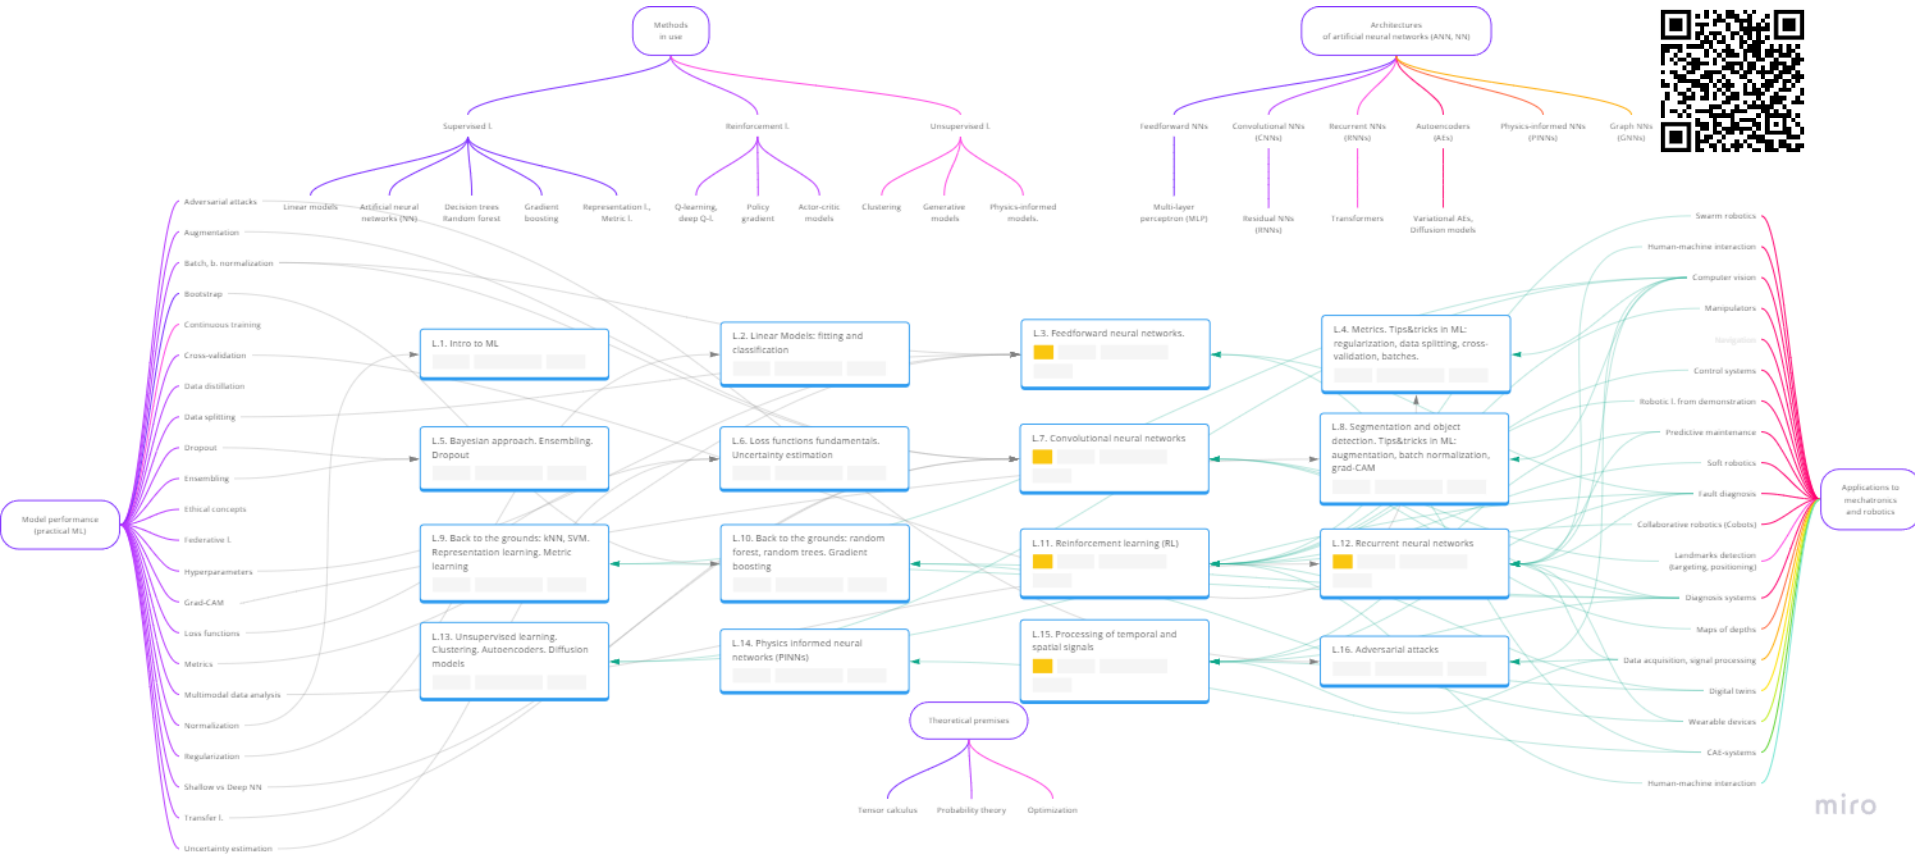
1. Explain how ML works: from basic ideas to real-world problems
2. Teach you how to build a model from scratch or use an open-source model to solve a problem
3. Help you take a few steps forward from educational problems to scientific ones

- [1] N. Akhtar and A. Mian. Threat of adversarial attacks on deep learning in computer vision: A survey. *Ieee Access*, 6:14410–14430, 2018.
- [2] M. Artemyev and A. Ashukha. Handbook on Machine Learning (in Russian). Yandex, 2024. URL <https://education.yandex.ru/handbook/ml>.
- [3] Y. Gal and Z. Ghahramani. Dropout as a bayesian approximation: Representing model uncertainty in deep learning. In *international conference on machine learning*, pages 1050–1059. PMLR, 2016.
- [4] I. Goodfellow, Y. Bengio, and A. Courville. *Deep Learning*. MIT Press, 2016. <http://www.deeplearningbook.org>.
- [5] G. E. Karniadakis, I. G. Kevrekidis, L. Lu, P. Perdikaris, S. Wang, and L. Yang. Physics-informed machine learning. *Nature Reviews Physics*, 3(6):422–440, 2021.
- [6] A. Kendall and Y. Gal. What uncertainties do we need in bayesian deep learning for computer vision? *Advances in neural information processing systems*, 30, 2017.
- [7] S. J. Prince. *Understanding Deep Learning*. The MIT Press, 2023. URL <http://udlbook.com>.
- [8] M. Raissi, P. Perdikaris, and G. E. Karniadakis. Physics-informed neural networks: A deep learning framework for solving forward and inverse problems involving nonlinear partial differential equations. *Journal of Computational physics*, 378:686–707, 2019.
- [9] R. R. Selvaraju, M. Cogswell, A. Das, R. Vedantam, D. Parikh, and D. Batra. Grad-cam: visual explanations from deep networks via gradient-based localization. *International journal of computer vision*, 128:336–359, 2020.



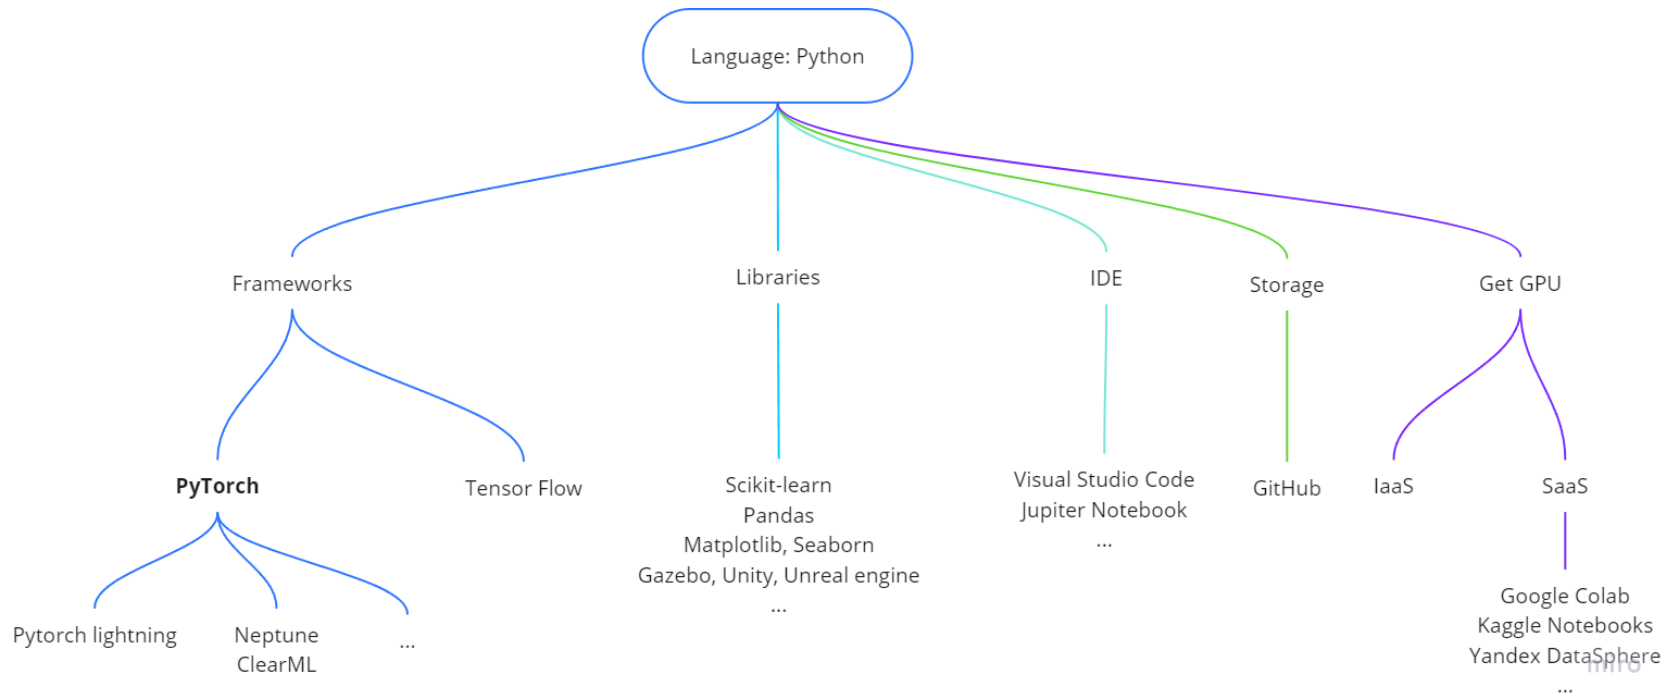
## Syllabus

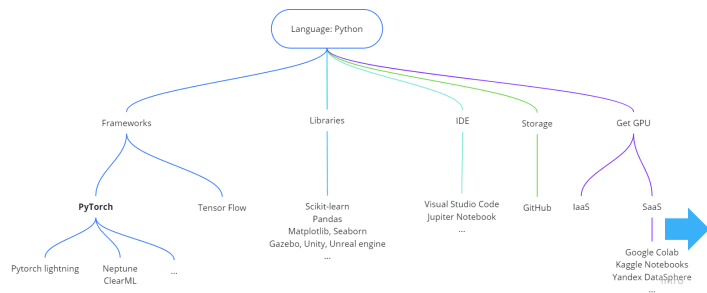
1. Intro to machine learning (ML) [2]
2. Linear models: fitting and classification [2]
3. Feedforward neural networks. Backpropagation [7, 2, 4]
4. Metrics. Tips and tricks in ML: regularization, data splitting, cross-validation, batches. [7, 2, 4]
5. Bayesian approach. Ensembling. Dropout [2, 3]
6. Loss functions fundamentals. Uncertainty estimation [7, 4, 2, 6]
7. Convolutional neural networks (CNNs). Residual neural networks [7, 2, 4]
8. Segmentation and object detection. Tips and tricks in ML: augmentation, batch normalization, grad-CAM [7, 2, 9]
9. Back to the grounds: kNN, SVM. Representation learning. Metric learning [2]
10. Back to the grounds: random forest, random trees. Gradient boosting [2]
11. Reinforcement learning (RL) [2, 7]
12. Recurrent neural networks (RNNs). Transformers [2, 7]
13. Unsupervised learning. Clustering. Autoencoders. Diffusion models [7, 2]
14. Physics informed neural networks (PINNs) [5, 8]
15. Processing of temporal and spatial signals [2]
16. Adversarial attacks [1]



## Course objectives

1. Explain how ML works: from basic ideas to real-world problems
2. Teach you how to build a model from scratch or use an open-source model to solve a problem
3. Help you take a few steps forward from educational problems to scientific ones





[Handbook on ML](#) by Yandex, 2022.

Критерий	Kaggle Notebooks	AWS SageMaker	Azure ML Studio	Google Colab	Yandex DataSphere
Вычислительные ресурсы	Ограничены	Масштабируемы, без серьезных ограничений	Масштабируемы, без серьезных ограничений	Ограничены	Масштабируемы, без ограничений времени использования
Ориентация	Соревнования по машинному обучению	Промышленная разработка, машинное обучение	Промышленная разработка, машинное обучение	Обучение, исследования	Промышленная разработка, машинное обучение
Интеграция с облачными хранилищами	Автоматическое монтирование датасетов	Создание датасетов из хранилищ AWS	Создание датасетов из хранилищ Azure	Упрощенный доступ к данным через Google Drive	Создание датасетов из хранилищ Yandex Cloud Object Storage, хранение датасетов в проекте
Поддержка коллективной работы	Обмен ноутбуками и датасетами с сообществом	Общие Docker-образы, унификация сред	Общие Docker-образы, унификация сред	Поддержка совместной работы через Google Drive	Коллективная работа над проектами и ресурсами
Интерфейс пользователя	Jupyter Notebook	Платформенно-ориентированный UI	Платформенно-ориентированный UI	Jupyter Notebook	Jupyter Notebook
Интеграция с Git	Доступна	Доступна	Доступна	Доступна через интерфейс командной строки	Доступна
Подход к данным и вычислительным ресурсам	Свансы с ограниченным временем использования	Отделение данных от вычислительных ресурсов	Отделение данных от вычислительных ресурсов	Сеансы с ограниченным временем использования, необходимость сохранения данных в Google Drive	Отделение данных от вычислительных ресурсов, возможность легкого переключения между конфигурациями VM
Особенности для образования и исследований	Соревнования по машинному обучению, обмен решениями	—	—	Широко используется в академических и образовательных целях	Оптимизация пользовательского опыта для студентов, интеграция с учебными курсами
Уникальные функции	Соревнования, обширное сообщество	Интегрированные решения для ML-разработки и развертывания	Визуальное конструирование ML-моделей	Простота использования, бесплатный доступ к ограниченным вычислительным ресурсам	DataSphere Jobs для удобной работы с вычислениями в командной строке, датасеты как виртуальные диски







## Books

[Handbook on Machine Learning](#) by M. Artemyev et al., Yandex, 2022 (in Russian)  
[Understanding Deep Learning](#) by Simon J.D. Prince, 2024  
[Practical Deep Learning / FastAI book](#) by Jeremy Howard  
[Deep Learning](#) by Ian Goodfellow and Yoshua Bengio and Aaron Courville, 2016.

## Online platforms, courses, resources

[Sirius](#) online courses on ML (in Russian)  
[Stepik](#) online courses (in Russian)  
[Hugging Face](#) online courses  
 Coursera is unavailable so far

[MIT Introduction to Deep Learning](#), MIT, 2024  
[Lecture Hall of the Faculty of Applied Mathematics and Informatics](#) (in Russian)  
[Fast AI](#), courses, software, book by Jeremy Howard  
[Deep Learning](#), course by Semyon Kozlov (in Russian), 2019  
  
[3Blue1Brown](#), Animated Math  
[PyTorch Tutorial](#) by Patrick Loeber, 2020

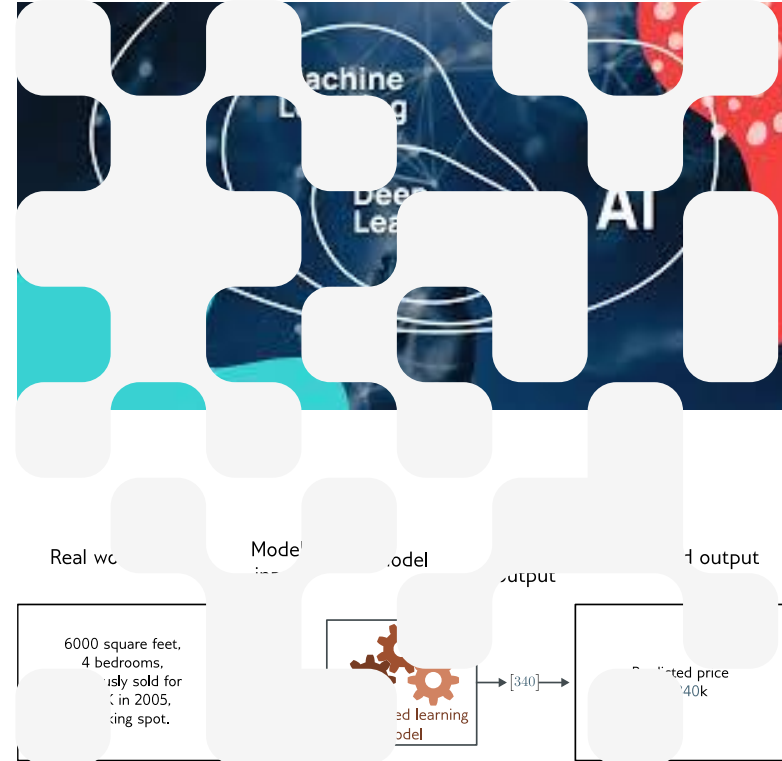
## #someLinks

Read here: <https://arxiv.org/>, <https://scholar.google.ru/>  
 Collect the references here: <https://mendeley.com/>  
 Draw here: <https://miro.com/app/dashboard/>  
 Write the text here: <https://www.overleaf.com/project>  
 Write the code here: <https://colab.research.google.com/>  
 Collect the code here: <https://github.com/>  
 Find the journal here: <https://journalfinder.elsevier.com/>  
 Find the conference here: <https://portal.core.edu.au/conf-ranks/?search=A>

# Agenda

- I. Logistics
- II. ML Overview:  $ML = E + T + P$
- III. Experience (E) in ML
- IV. ML Tasks (T)
- V. Performance measure (P) in ML

All models are wrong, but some are useful.  
/George Box/



$x \rightarrow \boxed{\phantom{x}} \rightarrow y$



# Deterministic vs Stochastic approaches to modeling



$$\rightarrow \frac{d^2x}{dt^2} = -g;$$




$$\rightarrow V(0) = V_0, \quad x(0) = x_0.$$

$$a = -g,$$

$$V = -gt + V_0,$$

$$x = -\frac{gt^2}{2} + V_0t + x_0.$$

Deterministic approach: solution of diff. eq. + initial conditions



②  $x = \varphi_0 + \varphi_1 t + \varphi_2 t^2$

③

Stochastic approach: experiment + approximation

$a = \varphi_0 + \varphi_1 t + \varphi_2 t^2$

## Terms

Искусственный интеллект /  
Artificial Intelligence (AI):

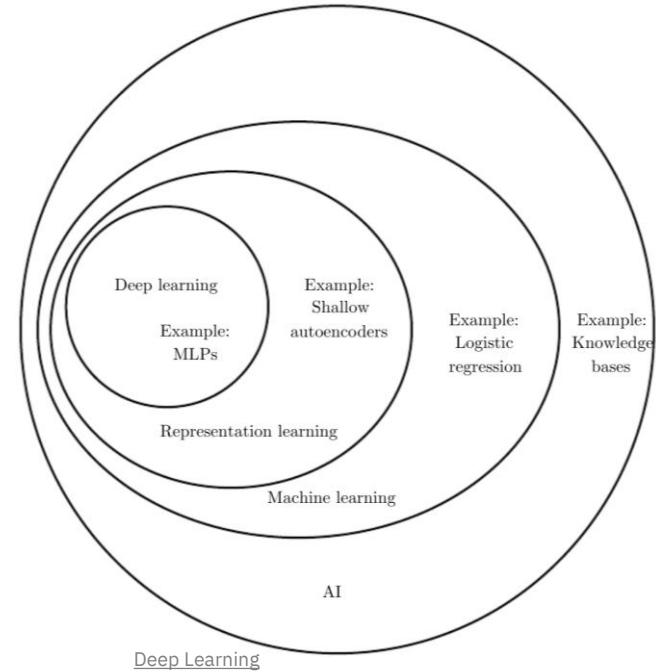
область информатики, занимающаяся моделированием разумного поведения в компьютерах / a branch of computer science dealing with the simulation of intelligent behavior in computers ([Merriam-Webster](#))

## AGI

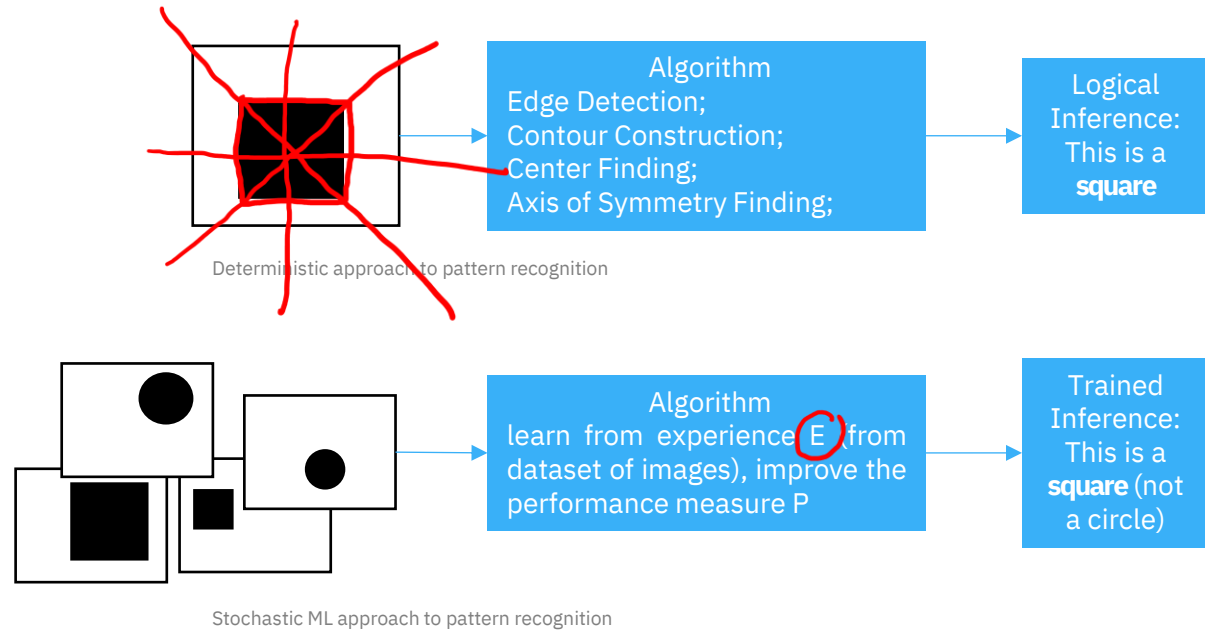
Машинное обучение / Machine Learning (ML):

область знаний, в которой компьютеры обучаются без явного программирования/ field of study that gives computers the ability to learn without being explicitly programmed (Arthur Samuel, 1959);

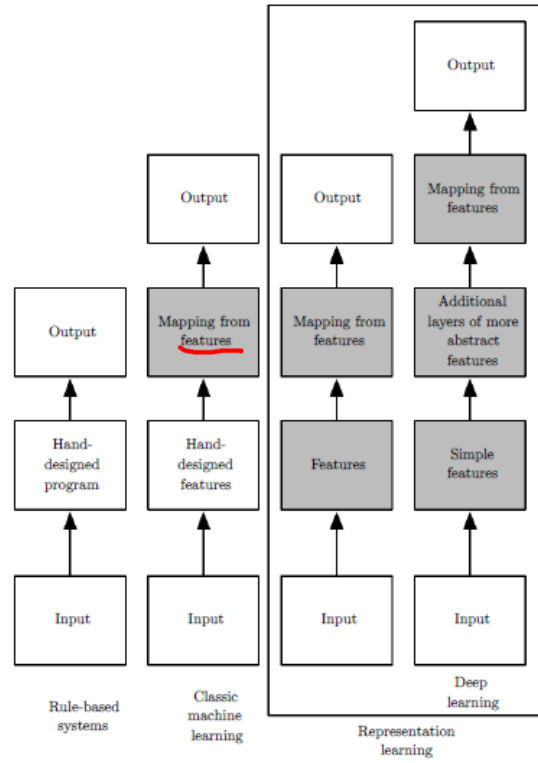
задача «З», в ходе решения которой программа обучается из опыта «О» и повышает меру качества «К» / well-posed learning problem: a computer program is said to learn from experience E with respect to some task T and some performance measure P, if its performance on T, as measured by P, improves with experience E (Tom Mitchell, 1998)




# Intuition



# Intuition

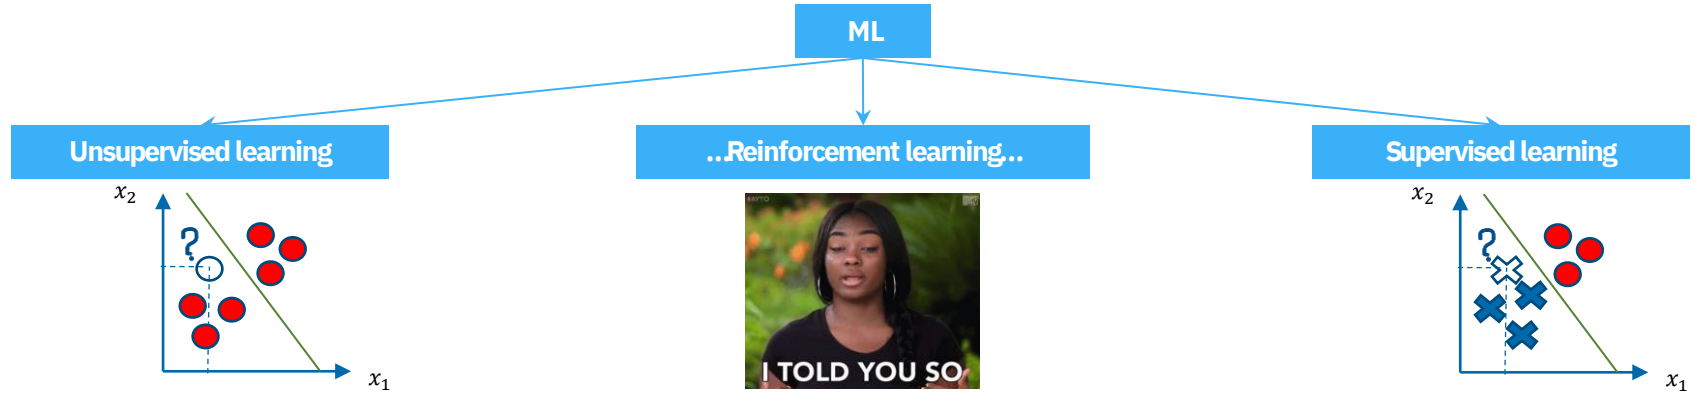


Deep Learning



Deep Learning

# Approaches to ML



# Prerequisites

## ML prerequisites

Линейная алгебра / Linear algebra


+

Мат. анализ / Calculus

+

Теор. вер. / Probability theory

Title	Geom. Analog	Notation in tensor / scalar forms	# of comp., in $\mathbb{R}^N$
Scalar		$a$	$N^0$
Vector		$\vec{a}, \mathbf{a}, [a_i], a_i$	$N^1$
Tensor (rank 2)		$\mathbf{T}_a, [a_{ij}], a_{ij}$	$N^2$



A *tensor* of rank  $n$  is a mathematical quantity characterized in  $N$ -dimensional space ( $\mathbb{R}^N$ ) by  $N^n$  components, each of which transforms according to a specific rule when the coordinate system is rotated\*:

$a' = a$  is for *scalar* (0-rank tensor);

$a'_i = \alpha_{ij} a_j$  is for *vector* (1-rank tensor),

$$(a'_1 = \alpha_{11}a_1 + \alpha_{12}a_2 + \alpha_{13}a_3,$$

$$a'_2 = \alpha_{21}a_1 + \alpha_{22}a_2 + \alpha_{23}a_3,$$

$$a'_3 = \alpha_{31}a_1 + \alpha_{32}a_2 + \alpha_{33}a_3);$$

$a'_{ik} = \alpha_{ij}\alpha_{km}a_{jm}$  is for *tensor* (2-rank tensor);

$a'_{i\dots k} = \alpha_{ij}\dots\alpha_{km}a_{j\dots m}$  is for *tensor* in general ( $n$ -rank tensor);

\* - The [Einstein summation notation](#) is used

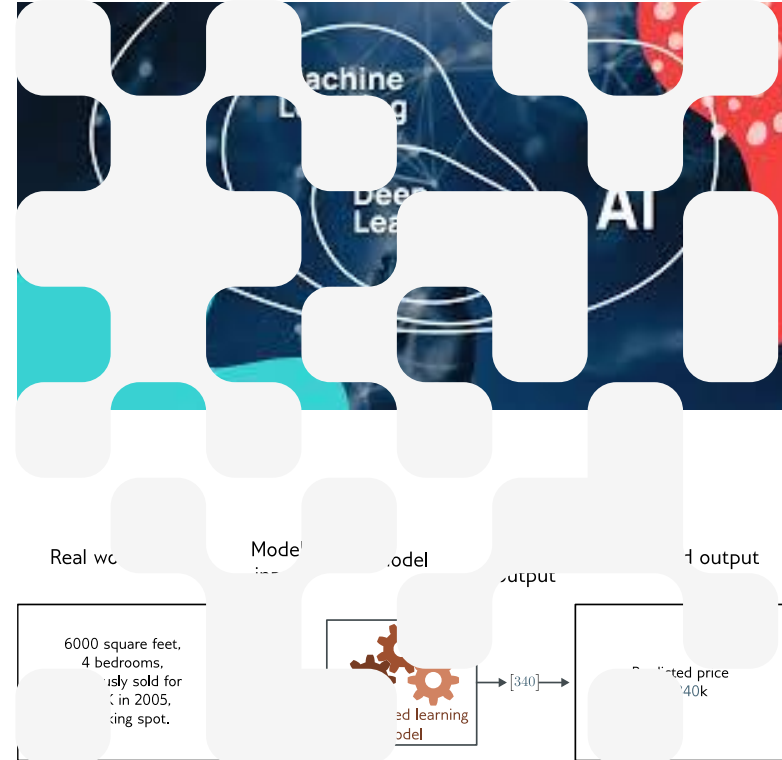




# Agenda

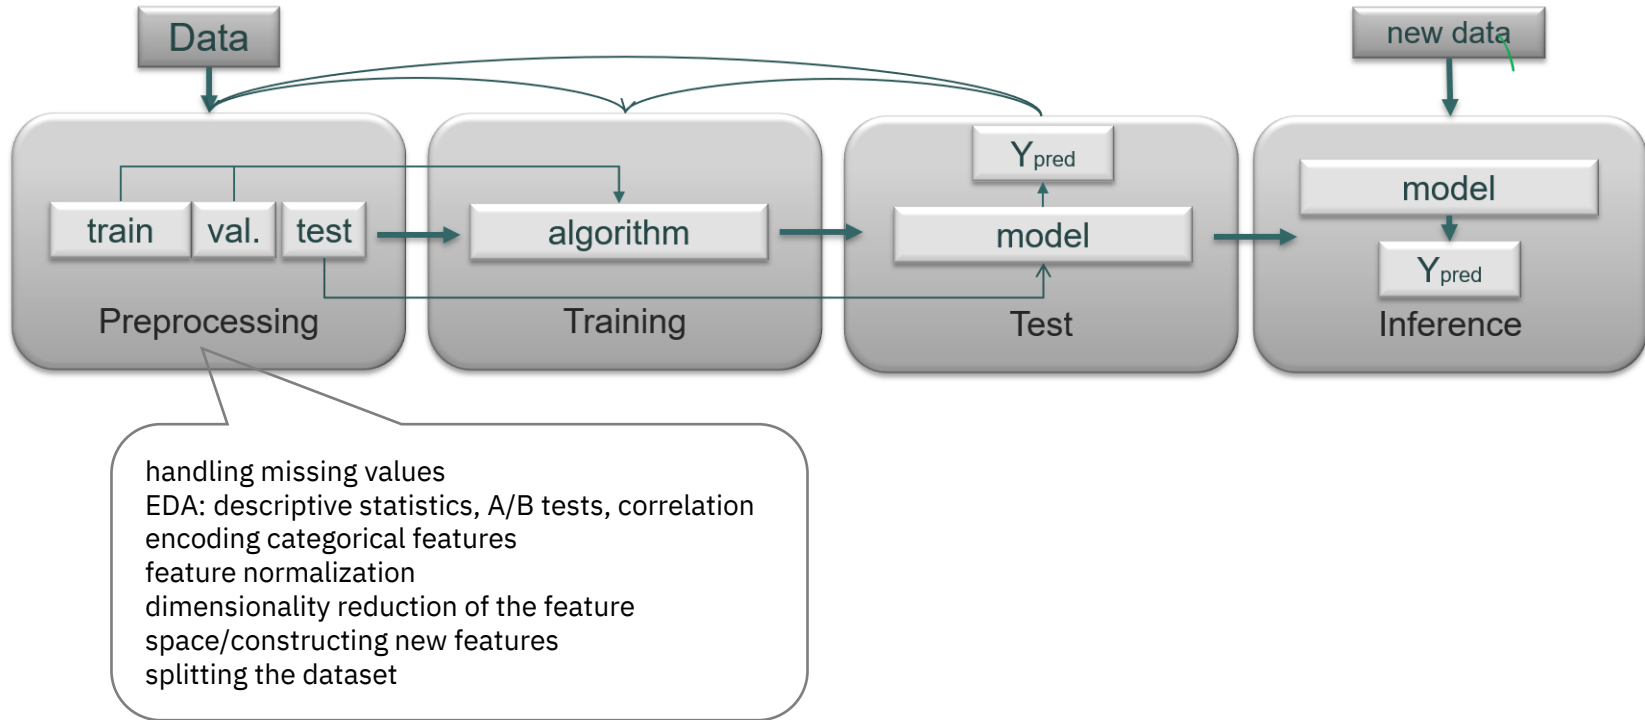
- I. Logistics
- II. ML Overview:  $ML = E + T + P$
- III. Experience (E) in ML
- IV. ML Tasks (T)
- V. Performance measure (P) in ML

All models are wrong, but some are useful.  
/George Box/





$$\vec{x} \rightarrow \boxed{\text{model}} \rightarrow y$$

## Flowchart for an ML model design (tabular data)






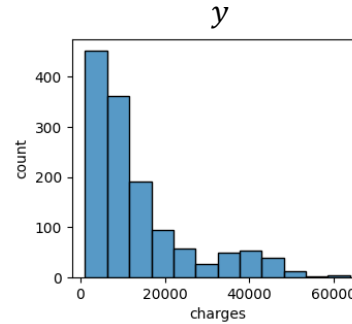


# Exploratory data analysis (EDA)

	age	sex	bmi	children	smoker	region	<u>charges</u>
0	19	female	27.900	0	yes	southwest	16884.92400
1	18	male	33.770	1	no	southeast	1725.55230
2	28	male	33.000	3	no	southeast	4449.46200
3	33	male	22.705	0	no	northwest	21984.47061
4	32	male	28.880	0	no	northwest	3866.85520



Medical insurance payout



# Encode categorical features

```
print(data.region.unique())
['southwest', 'southeast', 'northwest', 'northeast']
```

```
pd.get_dummies(data[['region']], dtype=int)
```

region	region_northeast	region_northwest	region_southeast	region_southwest
southwest	0	0	0	1
southeast	0	0	1	0
southeast	0	0	1	0
northwest	0	1	0	0
northwest	0	1	0	0

⚠ warning:  $x_4 = 1 - \sum_{i=1}^3 x_i$


```
pd.get_dummies(data[['region']], dtype=int,
drop_first=True)
```

region_northwest	region_southeast	region_southwest
0	0	1
0	1	0
0	1	0
1	0	0
1	0	0

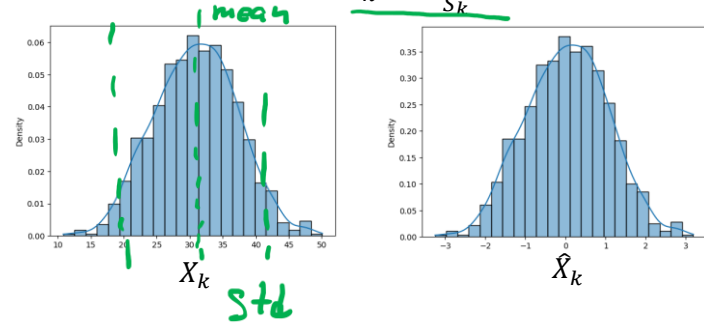
Medical insurance payout

# Normalize features

1. MinMaxScaler:  $\hat{X}_k = \frac{X_k - x_{kmin}}{x_{kmax} - x_{kmin}}$



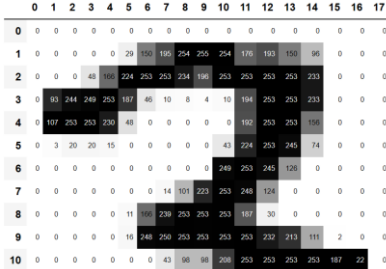
2. StandardScaler:  $\hat{X}_k = \frac{X_k - \bar{x}_k}{S_k}$



```
from sklearn.preprocessing import MinMaxScaler
scaler = MinMaxScaler()
X_num_norm = scaler.fit_transform(data[['age', 'bmi', 'children']])
```

	sex_male	smoker_yes	region_northwest	region_southeast	region_southwest	age_norm	bmi_norm	children_norm	ln_charges
0	0.0	1.0	0.0	0.0	1.0	0.021739	0.321227	0.0	9.734236
1	1.0	0.0	0.0	1.0	0.0	0.000000	0.479150	0.2	7.453882
2	1.0	0.0	0.0	1.0	0.0	0.217391	0.458434	0.6	8.400763
3	1.0	0.0	1.0	0.0	0.0	0.326087	0.181464	0.0	9.998137
4	1.0	0.0	1.0	0.0	0.0	0.304348	0.347592	0.0	8.260455

# Data: input(s) is (are) stationary

Numeric			Visual	Textual
Feature(s)			Image	Word
Feature	Value	Descript.		<p>“cat”</p>
x_0	2	# rooms		
...				
x_m	5.5	Distance		

$$X = [x_i] = \begin{bmatrix} x_0 \\ \dots \\ x_m \end{bmatrix}$$

Architecture:

Perceptron,  
decision tree,  
1D-CNN,  
Transformer

$$X = [x_{ij}] = \begin{bmatrix} x_{01} & \dots & x_{0n} \\ \dots & \dots & \dots \\ x_{0m} & \dots & x_{mn} \end{bmatrix}$$

Convolutional  
Neural Network  
(CNN),  
Transformer

$$X = [x_i] = \begin{bmatrix} x_0 \\ \dots \\ x_m \end{bmatrix}$$

Transformer

# Data: input(s) is (are) non-stationary

$$\vec{a} \cdot \vec{b} = \|\vec{a}\| \|\vec{b}\| \cos(\angle \vec{a} \vec{b})$$

↓  
sim

Numeric	Visual	Textual
---------	--------	---------


Sequence of  
features

Sequence of  
images


Sequence of  
words

"The|cat's|eyes|  
sparkled|with|  
..."  
→ tokens

→ embeddings  $\mathbb{R}^N$



Transformer



Architecture:



Recurrent  
network (RN),  
Transformer

CNN + LSTM,  
Transformer





Data: output(s) is (are) stationary

Numeric			Visual	Textual
Number(s)			Image	Word
Feature	...	Price, MRub		"curiosity"
x_0	...	10		
...	...			
x_m	...			

$y$

$$Y = [y_{ij}] = \begin{bmatrix} y_{01} & \dots & y_{0n} \\ \dots & \dots & \dots \\ y_{0m} & \dots & y_{mn} \end{bmatrix}$$

$$Y = [y_i] = \begin{bmatrix} y_0 \\ \dots \\ y_m \end{bmatrix}$$

Architecture:

Perceptron,  
decision tree,  
1D-CNN,  
Transformer


Convolutional  
Neural Network  
(CNN)

Transformer


# Agenda

- I. Logistics
- II. ML Overview:  $ML = E + T + P$
- III. Experience (E) in ML
- IV. ML Tasks (T)
- V. Performance measure (P) in ML

All models are wrong, but some are useful.  
 — George Box



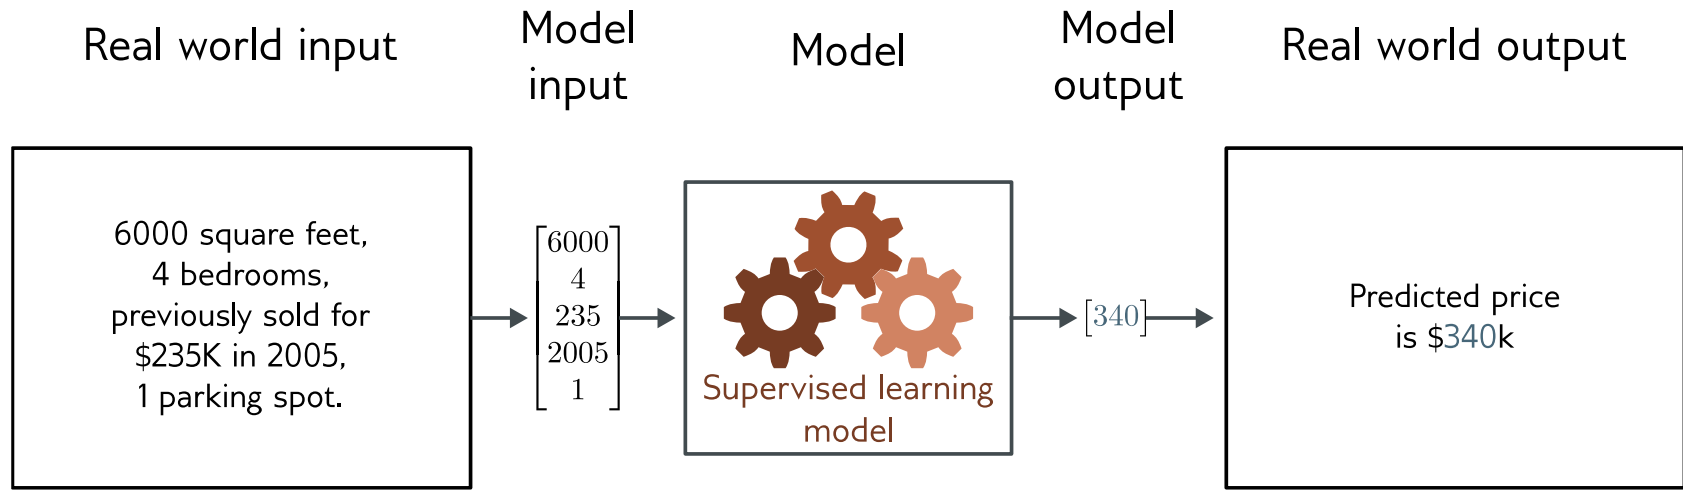
## An ML model intuition



### Supervised learning intuition

S. J. Prince. Understanding Deep Learning. MIT Press, 2023. URL <http://udlbook.com>.

# Regression

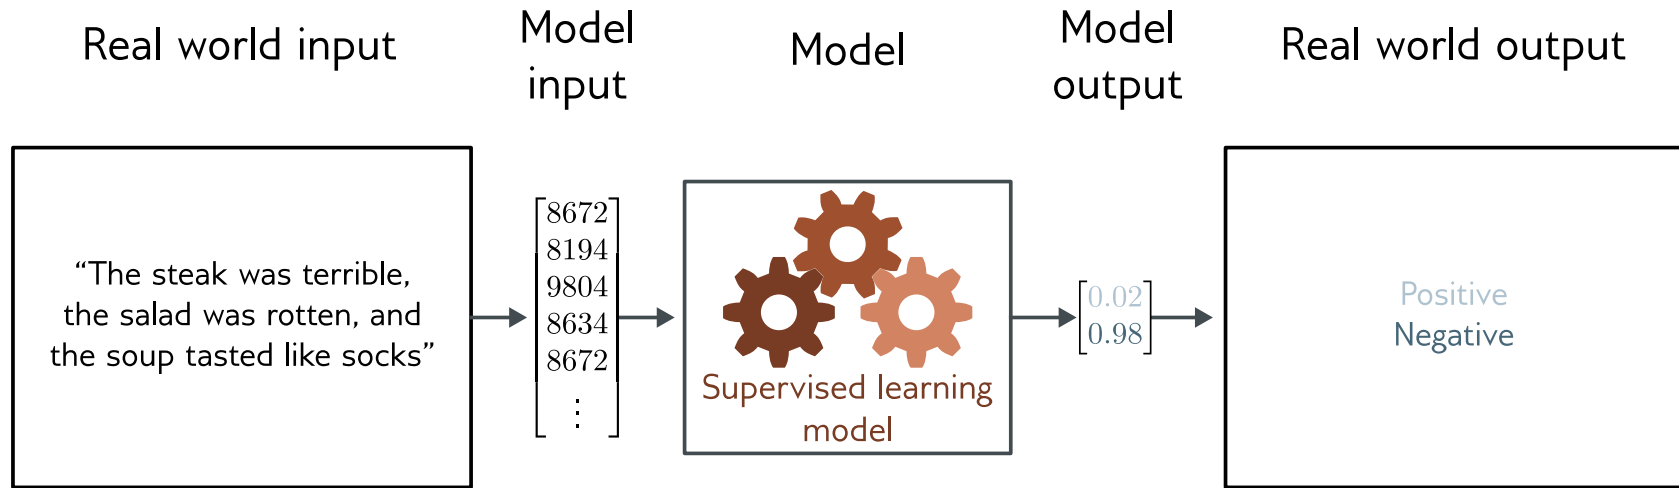


Univariate regression problem (one output, real value)  
Fully connected network

## Supervised learning intuition

S. J. Prince. Understanding Deep Learning. MIT Press, 2023. URL <http://udlbook.com>.

## Text classification

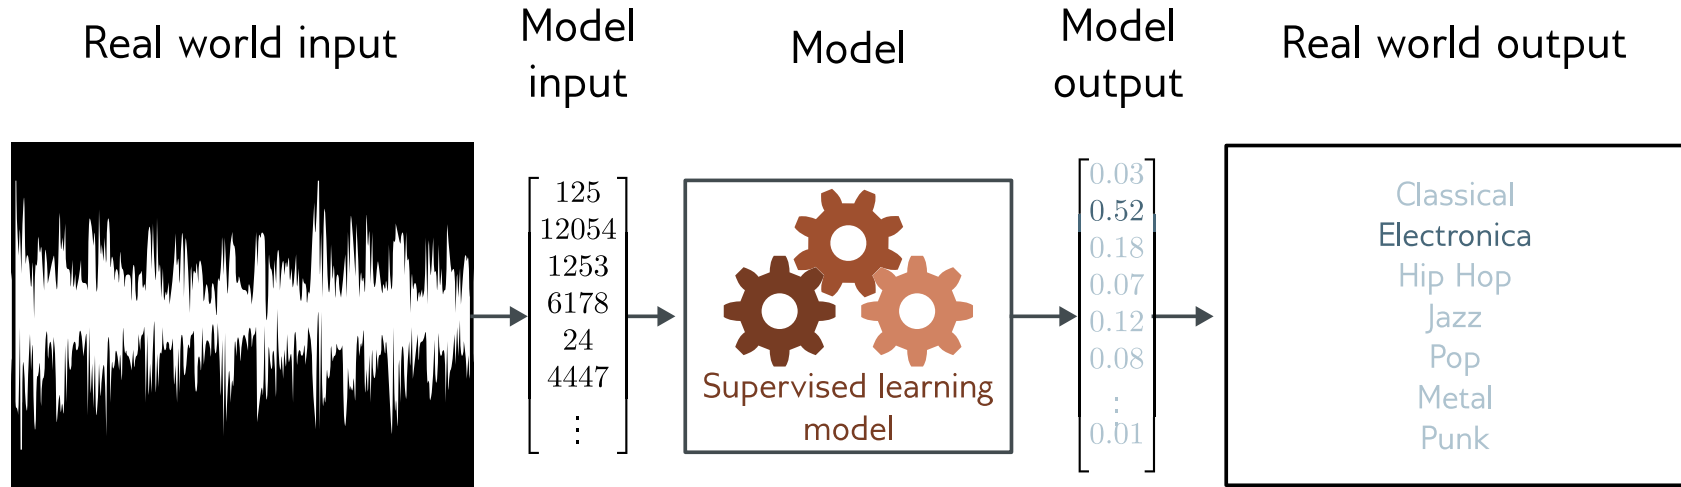


Binary classification problem (two discrete classes)  
Transformer network

### Supervised learning intuition

S. J. Prince. Understanding Deep Learning. MIT Press, 2023. URL <http://udlbook.com>.

## Music genre classification

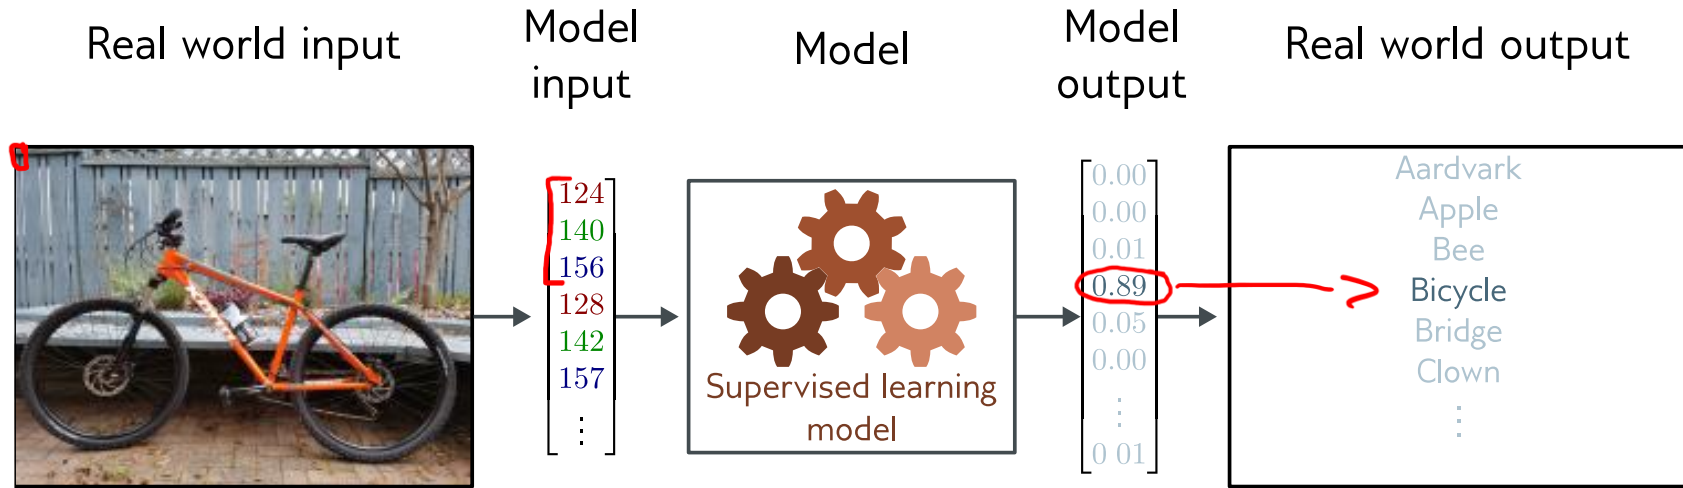


Multiclass classification problem (discrete classes, >2 possible values)  
Recurrent neural network (RNN), Transformers

### Supervised learning intuition

S. J. Prince. Understanding Deep Learning. MIT Press, 2023. URL <http://udlbook.com>.

# Image classification




Multiclass classification problem (discrete classes, >2 possible classes)

Convolutional networks, Transformers

## Supervised learning intuition

S. J. Prince. Understanding Deep Learning. MIT Press, 2023. URL <http://udlbook.com>.

# Image segmentation



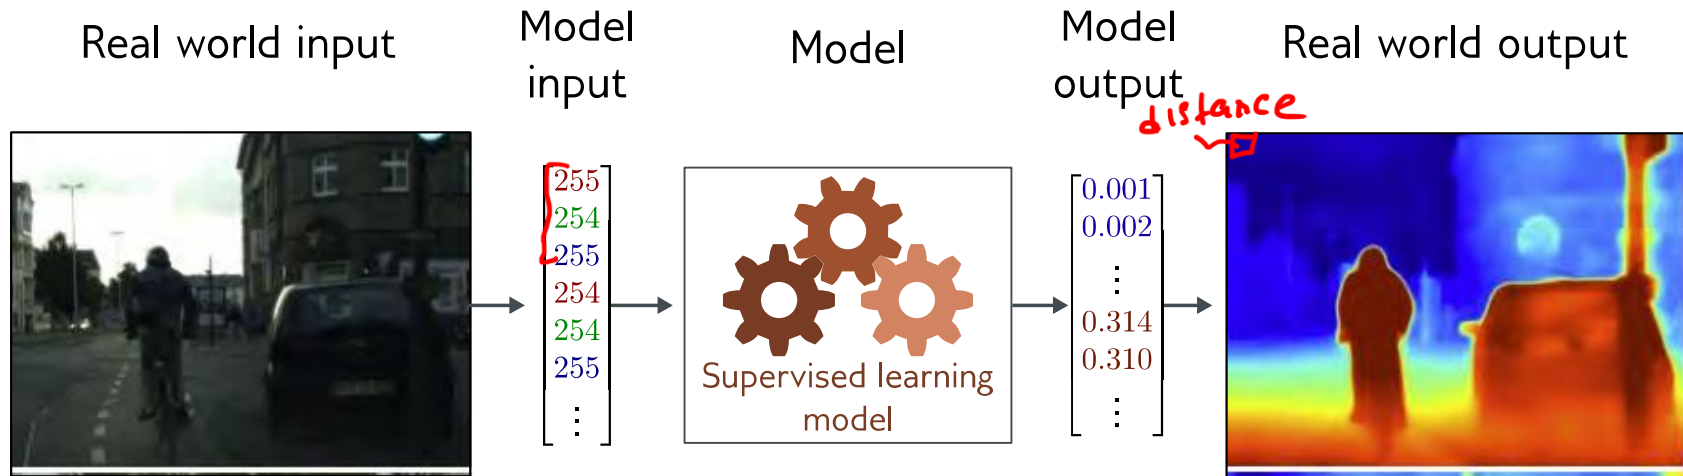
Multivariate binary classification problem (many outputs, two discrete classes)  
 Convolutional encoder-decoder network

## Supervised learning intuition

S. J. Prince. Understanding Deep Learning. MIT Press, 2023. URL <http://udlbook.com>.



# Depth estimation

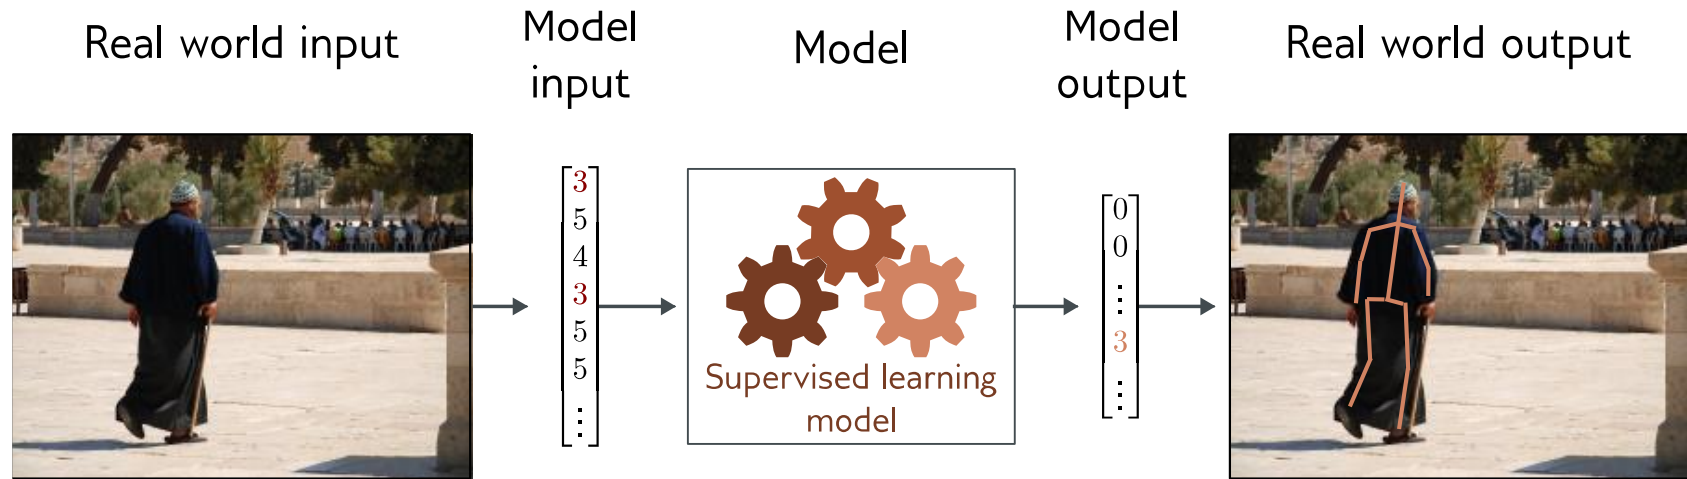


Multivariate regression problem (many outputs, continuous)  
 Convolutional encoder-decoder network

## Supervised learning intuition

S. J. Prince. Understanding Deep Learning. MIT Press, 2023. URL <http://udlbook.com>.

## Pose estimation



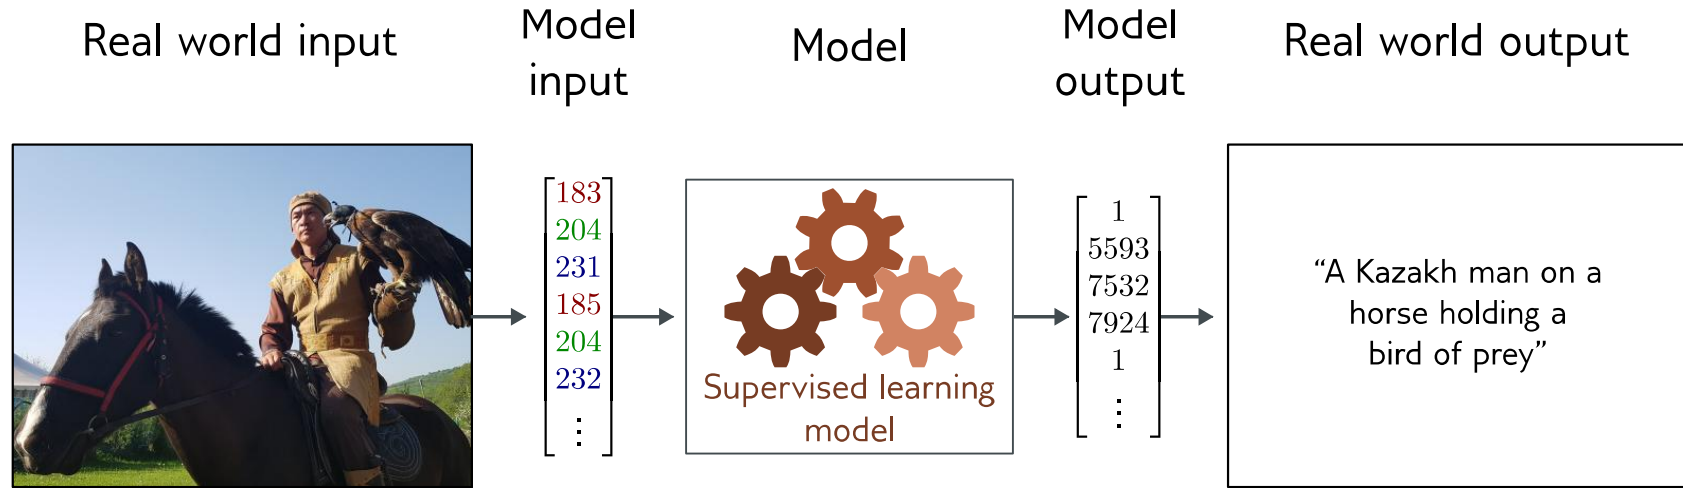
Multivariate regression problem (many outputs, continuous)

Convolutional encoder-decoder network

### Supervised learning intuition

S. J. Prince. Understanding Deep Learning. MIT Press, 2023. URL <http://udlbook.com>.

# Image captioning




Visual data analysis and text generation problem  
Transformer networks

## Supervised learning intuition

S. J. Prince. Understanding Deep Learning. MIT Press, 2023. URL <http://udlbook.com>.

# Games



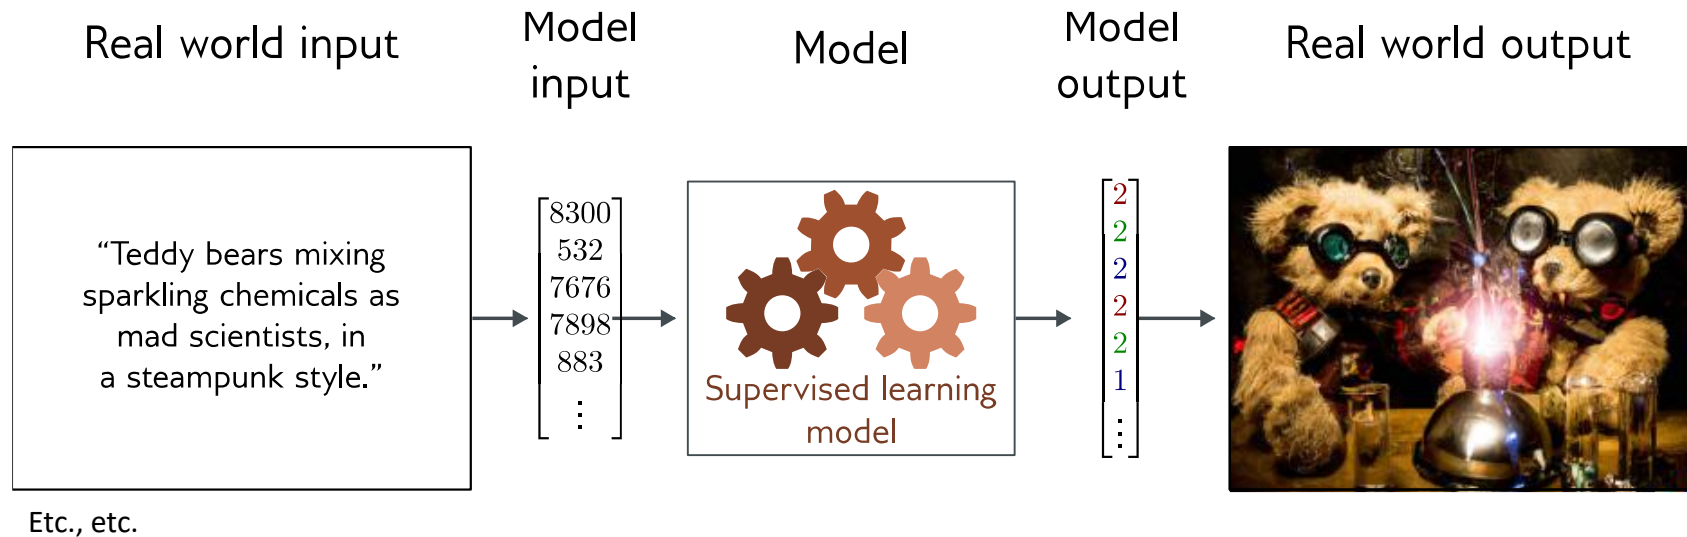
Decision making problem

Feedforward networks, CNNs, transformer networks

## Supervised learning intuition

S. J. Prince. Understanding Deep Learning. MIT Press, 2023. URL <http://udlbook.com>.

## Image generation from text



### Supervised learning intuition

S. J. Prince. Understanding Deep Learning. MIT Press, 2023. URL <http://udlbook.com>.


# Agenda

- I. Logistics
- II. ML Overview:  $ML = E + T + P$
- III. Experience (E) in ML
- IV. ML Tasks (T)
- V. Performance measure (P) in ML

All models are wrong, but some are useful.  
/George Box/






# Performance Measure intuition



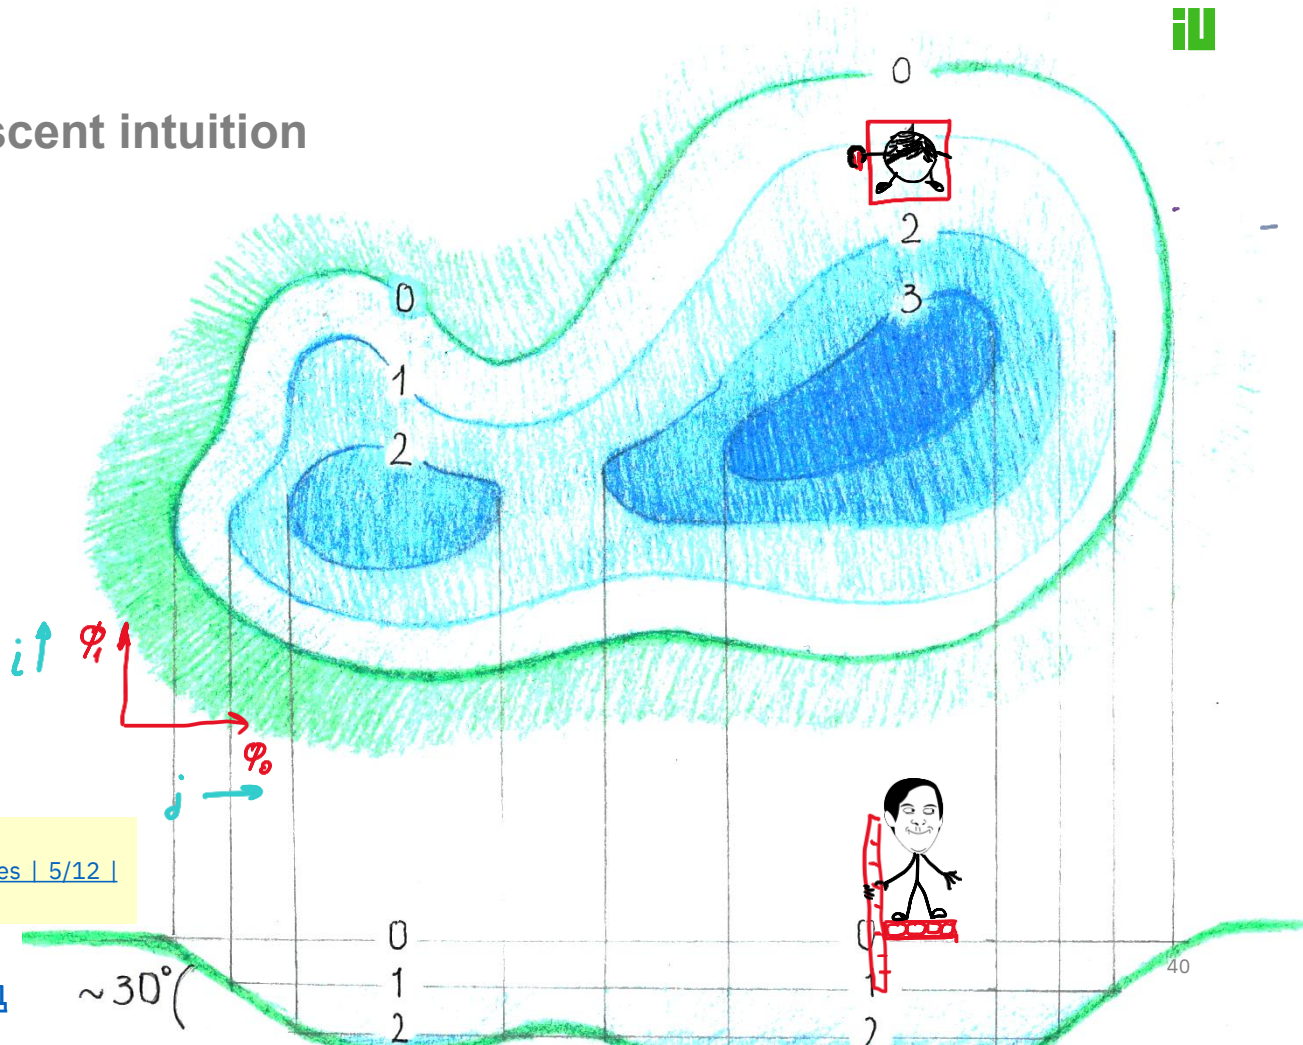

$$L_R = \frac{1}{2m} \sum_{i=1}^m \left( y^{(i)} - h^{(i)} \right)^2$$

Model predicts output  $h$  given input  $x$



Supervised learning intuition: S. J. Prince. Understanding Deep Learning. MIT Press, 2023. URL <http://udlbook.com>.

# Gradient ascent/descent intuition



Want to know more?


[DeepMind x UCL | Deep Learning Lectures | 5/12 | Optimization for Machine Learning](#)

Пруд

$\sim 30^\circ$





# Performance Measure intuition



$$L_R = \frac{1}{2m} \sum_{i=1}^m \left( y^{(i)} - h^{(i)} \right)^2$$

Model predicts output  $h$  given input  $x$



Supervised learning intuition: S. J. Prince. Understanding Deep Learning. MIT Press, 2023. URL <http://ud1book.com>.

## Think about it



1. Why it is not recommended to differentiate functions obtained with approximation?
2. How can we estimate that *data* is enough to solve a problem with an ML model?
3. Why small noise in *data* may be usefull?
4. Is it possible to make an AI model that solves all problems at once?
5. Is it possible to solve a problem using a complex *performance measure*? E.g. minimize something and maximize something else simultaneously?

Thank you for your attention!

[a.kornaev@innopolis.ru](mailto:a.kornaev@innopolis.ru), @avkornaev



