

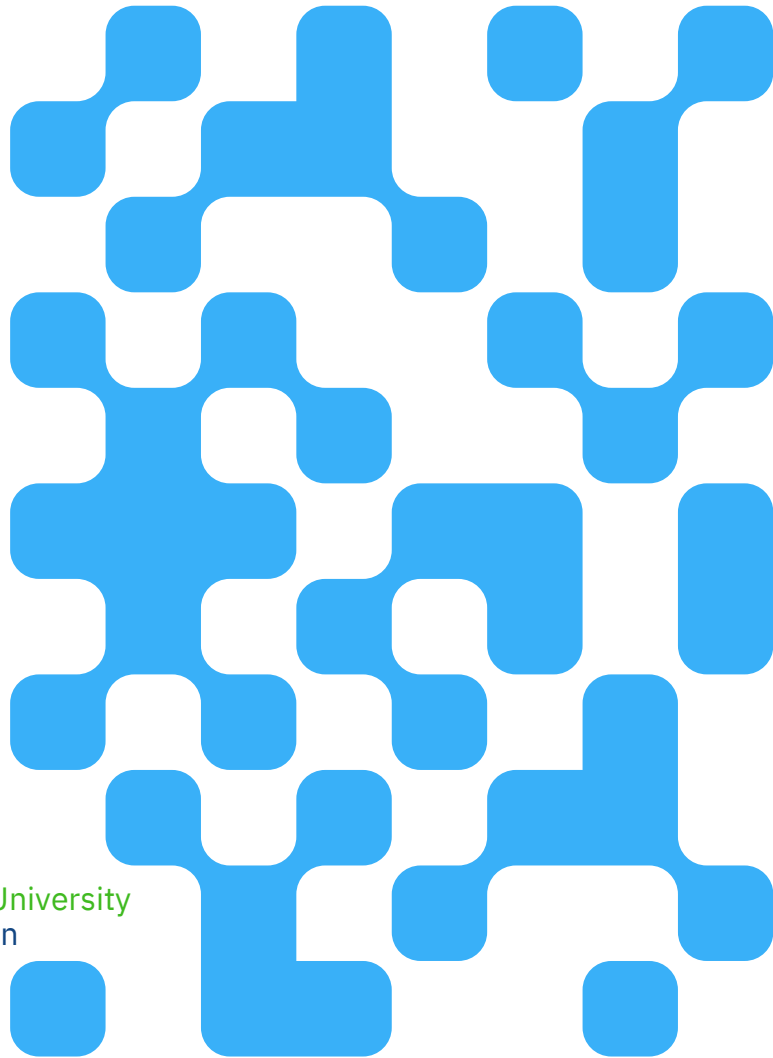


# Machine Learning

2025 (ML-2025)

Lecture 3. Metrics and Metric Approach

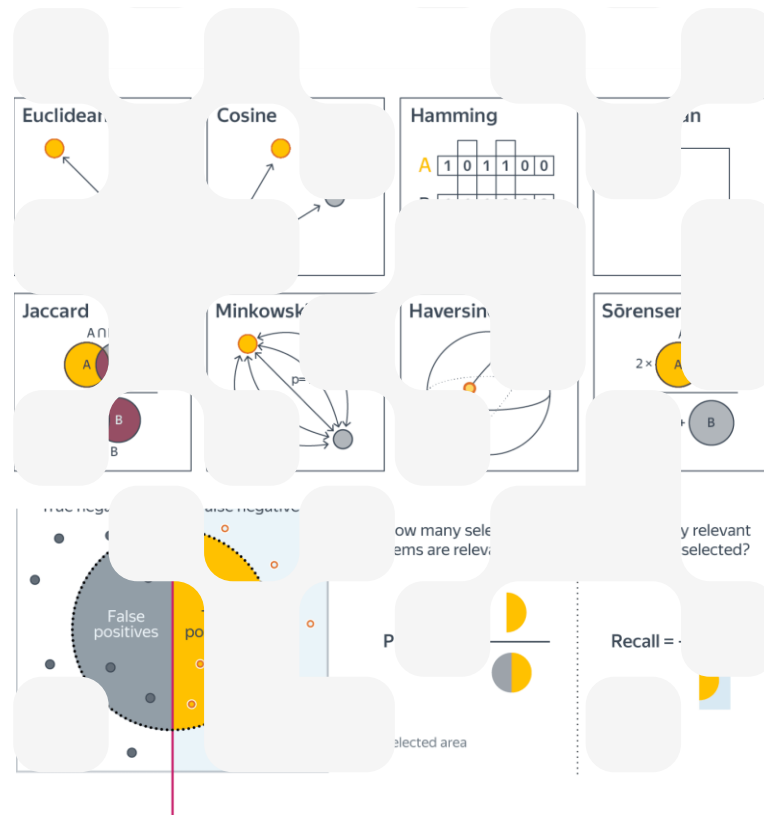
by Alexei Kornaev, Dr. Sc., Assoc. Prof., Robotics and CV, [Innopolis University](#)  
Researcher at the RC for AI, [National RC for Oncology n.a. NN Blokhin](#)



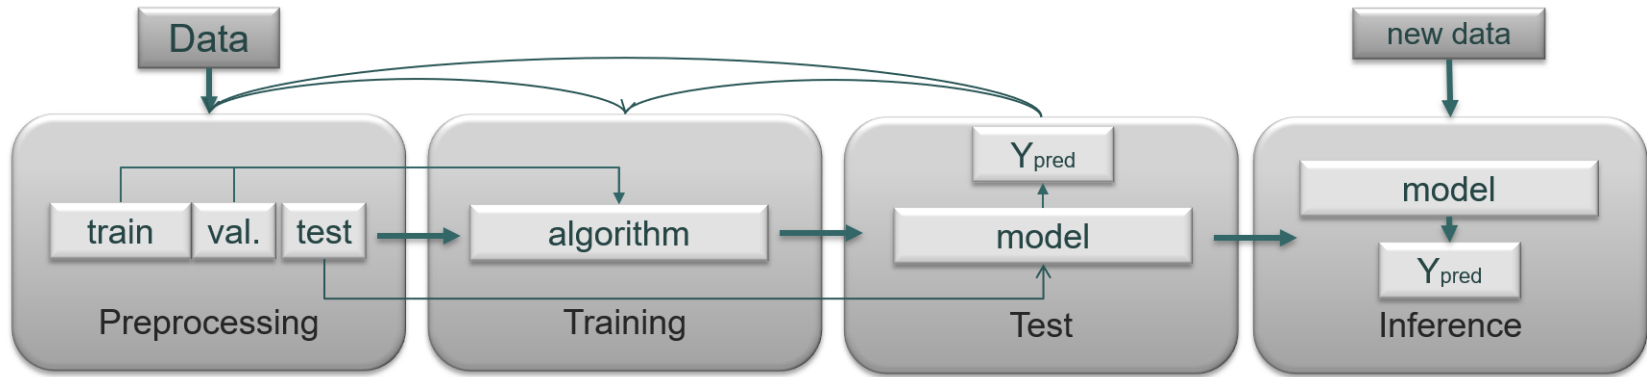
# Agenda

- I. Quality metrics in ML
- II. Data splitting, cross-validation
- III. Regularization
- IV. Batches
- V. Metric approach in ML: k-nearest neighbors

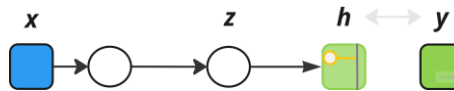
All models are wrong, but some are useful.  
/George Box/



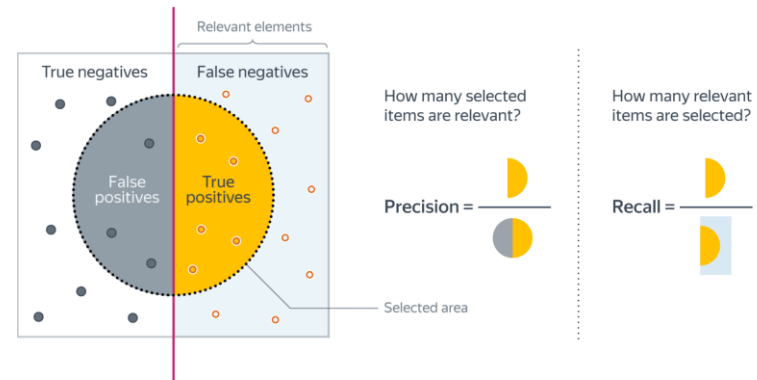
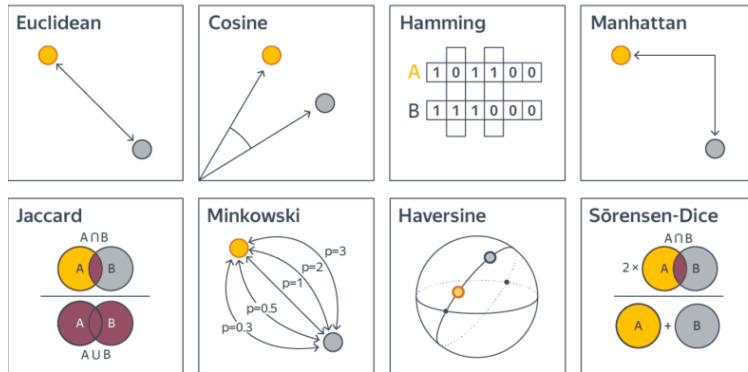
## Flowchart for an ML model design



**Metrics** can be used both in learning (metric learning) and in estimating a model's quality



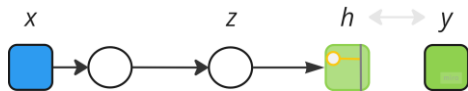
Model predicts output  $h$  given input  $x$



[Metrics from the Yandex Handbook](#)

[Classification metrics from the Yandex Handbook](#)

# Metrics that estimate model's quality. Binary classification



Model predicts output  $h$  given input  $x$

1. Accuracy:  $acc(\mathbf{h}, \mathbf{y}) = \frac{1}{m} \sum_{i=1}^m (h^{(i)} == y^{(i)})$ , or (and) error rate:  $error\ rate = 1 - acc(\mathbf{h}, \mathbf{y})$

$$acc(\mathbf{h}, \mathbf{y}) = \frac{TP+TN}{TP+TN+FP+FN}$$

$acc =$

Predicted	True
1	0
0	0
1	1
0	0
1	0
0	0
0	0
1	1
0	0
0	0

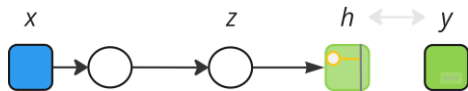
$acc =$

Predicted	True
0	1
0	0
0	0
0	0
0	1
0	0
0	0
0	0
0	0
0	0

Predicted class		
Positive	Negative	
TP	FN	Positive
FP	TN	Negative
		True class

[Confusion matrix](#)

# Metrics that estimate model's quality. Binary classification



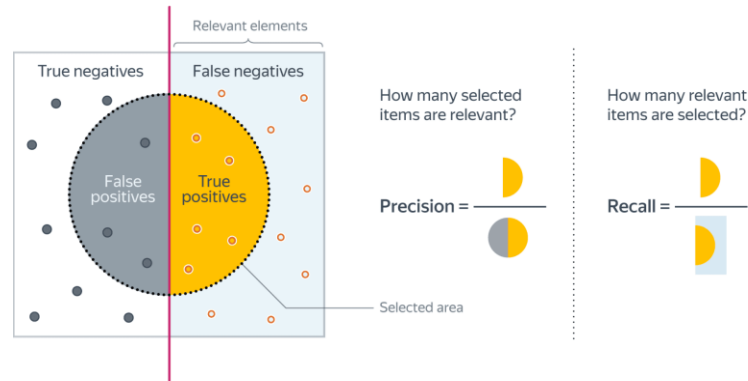
Model predicts output  $h$  given input  $x$

$$2. \text{Precision} = \frac{TP}{TP+FP}, \text{Recall} = \frac{TP}{TP+FN}, F_1 = \frac{\text{Recall} \cdot \text{Precision}}{\text{Recall} + \text{Precision}}.$$

Predicted class		
Positive	Negative	
TP	FN	Positive
FP	TN	Negative
		True class

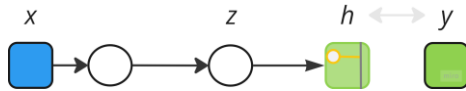
Predicted	True
1	0
0	0
1	1
0	0
1	0
0	0
0	0
1	1
0	0
0	0

Predicted	True
0	1
0	0
0	0
0	0
0	1
0	0
0	0
0	0
0	0
0	0



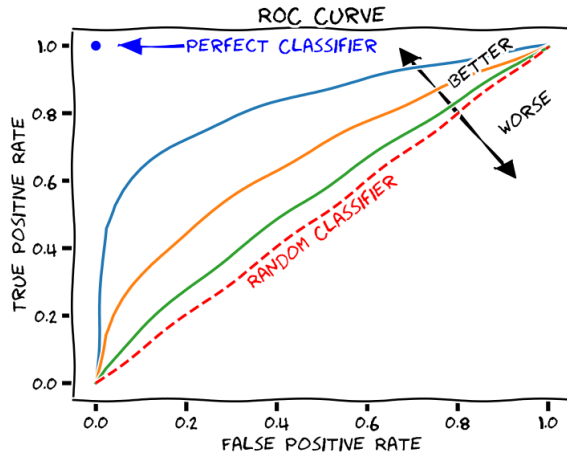
[Precision, recall, F1-score intuition](#)

# Metrics that estimate model's quality. Binary classification

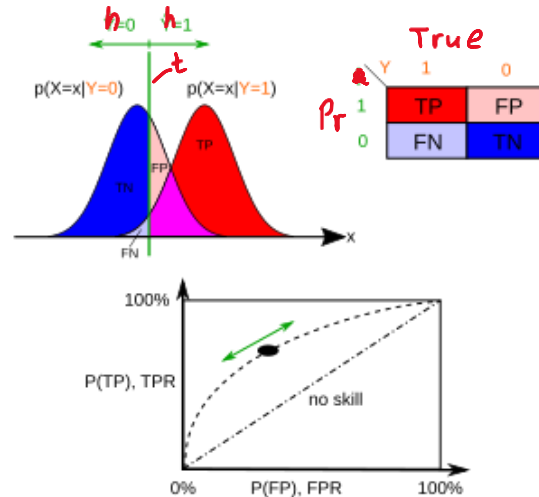


Model predicts output  $h$  given input  $x$

1. True positive rate  $TPR = Recall = \frac{TP}{TP+FN}$ , false positive rate  $FPR = \frac{FP}{FP+TN}$ , Receiver operating characteristic (ROC), area under curve  $AUC$ .



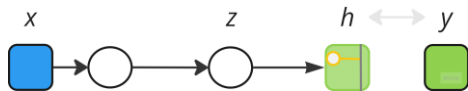
[ROC curve intuition](#)



Overall, the optimization of precision and recall proceeds as follows:

- train the model on a loss function;
- obtain metric graphs depending on the threshold using real predictions on the validation set, by iterating over different thresholds from 0 to 1;
- select the desired combination of precision and recall.

# Metrics that estimate model's quality. Fitting



Model predicts output  $h$  given input  $x$

1. Mean squared error (MSE):

$$MSE = \frac{1}{m} \sum_{i=1}^m (h^{(i)} - y^{(i)})^2;$$

2. Mean absolute error (MAE):

$$MAE = \frac{1}{m} \sum_{i=1}^m |h^{(i)} - y^{(i)}|;$$

3. Root MSE (RMSE):

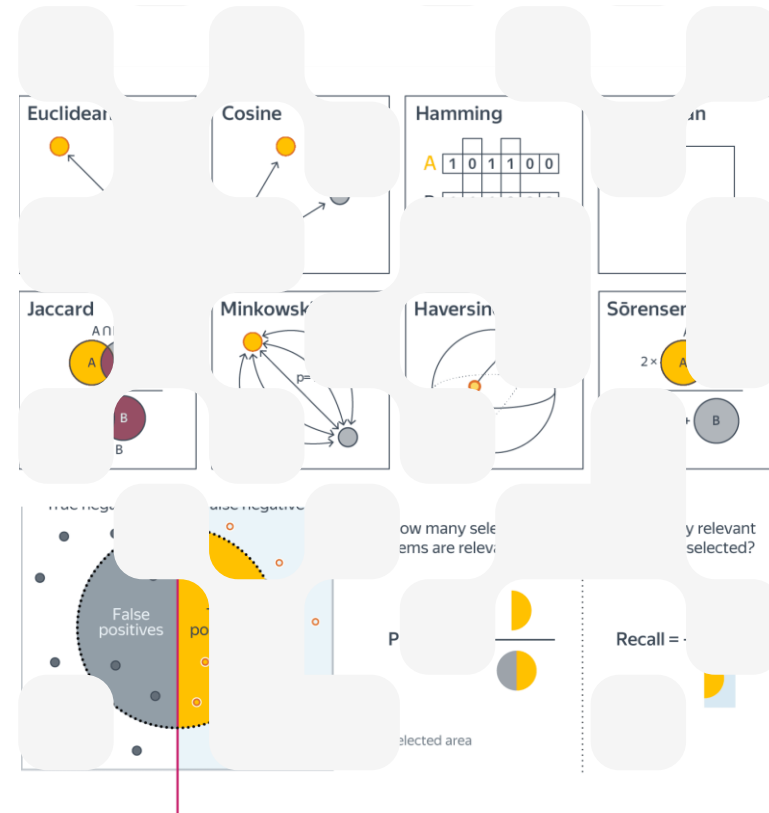
$$RMSE = \sqrt{\frac{1}{m} \sum_{i=1}^m (h^{(i)} - y^{(i)})^2}.$$



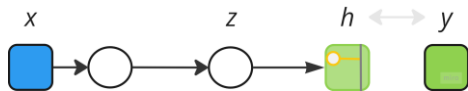
# Agenda

- I. Quality metrics in ML
- II. Data splitting, cross-validation
- III. Regularization
- IV. Batches
- V. Metric approach in ML: k-nearest neighbors

All models are wrong, but some are useful.  
/George Box/



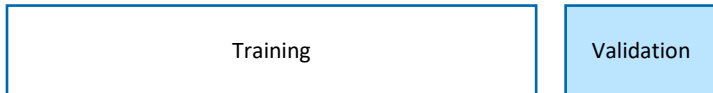
# Overfitting intuition and data splitting



Model predicts output  $h$  given input  $x$

Model *parameters* are determined during the solution of the ML problem, e.g. model weights. *Hyperparameters* are set by the user, usually not in a single way, and their values affect the values of the sought parameters.

Training data splitting by default: training set, validation set



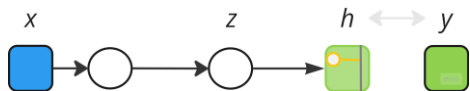
K-Fold splitting (normally used when the dataset is small)



Algorithm. The dataset is divided into  $k$  equal parts. Next,  $k$  iterations occur, during each of which one fold serves as the validation set, and the union of the remaining folds serves as the training set. The model is trained on  $k-1$  folds and tested on the remaining one. The final score of the model is obtained either by averaging the resulting test results or by measuring it on a held-out test set that did not participate in cross-validation.

Testing procedure: the accuracy of the selected trained model is checked on a new (test set)!!!

# Overfitting intuition and data splitting



Model predicts output  $h$  given input  $x$

$$h = \theta_j x^j, (j = 0, \dots, d)$$

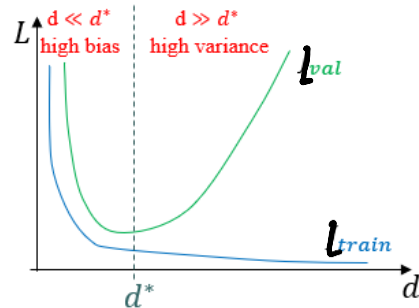
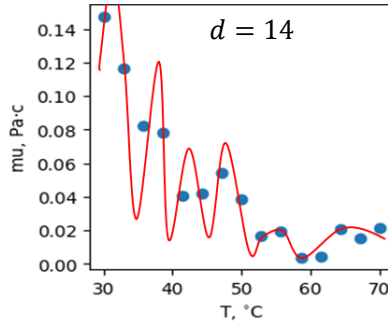
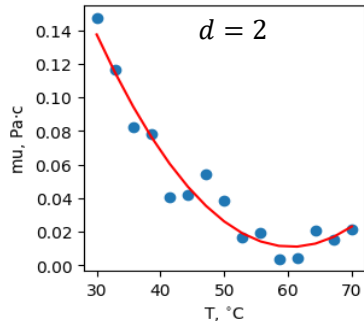
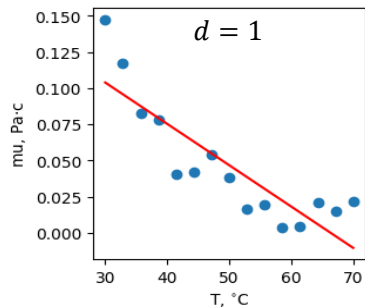
Model *parameters* are determined during the solution of the ML problem, e.g. model weights. *Hyperparameters* are set by the user, usually not in a single way, and their values affect the values of the sought parameters.

Training  $\{(x_i, y_i)\}$

validation

test

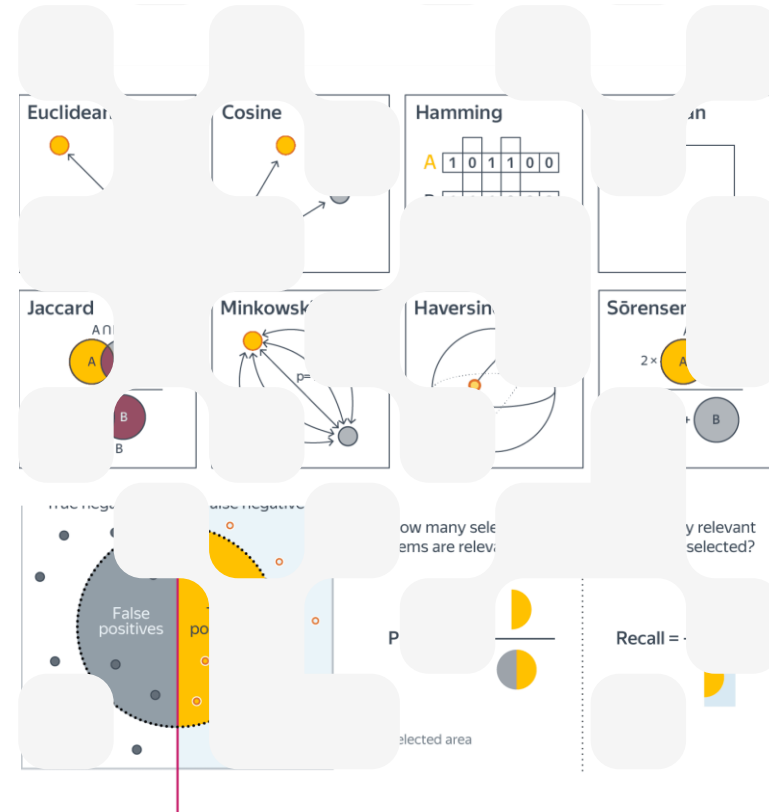
$$L = \frac{1}{2m} \left[ \sum_{i=1}^m (h^{(i)} - y^{(i)})^2 \right] \Rightarrow \min.$$



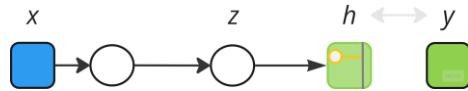
# Agenda

- I. Quality metrics in ML
- II. Data splitting, cross-validation
- III. Regularization
- IV. Batches
- V. Metric approach in ML: k-nearest neighbors

All models are wrong, but some are useful.  
/George Box/



# Overfitting intuition and regularization (L2)



Model predicts output  $h$  given input  $x$

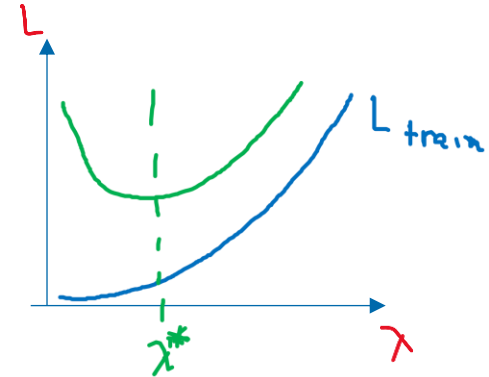
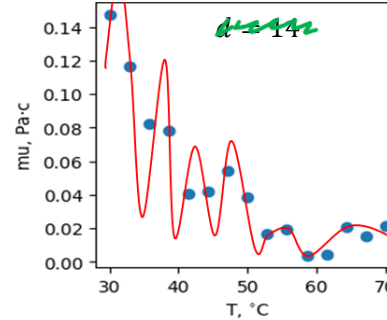
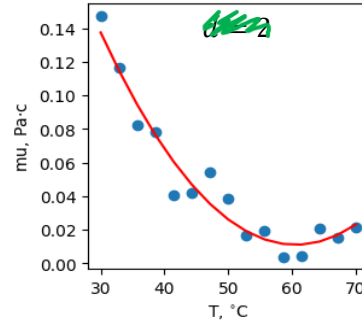
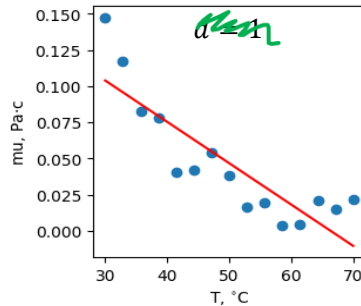
Model *parameters* are determined during the solution of the ML problem. For example, in regression problems, the parameters are the components of the matrix of weights  $\phi$ . *Hyperparameters* are set by the user, usually not in a single way, and their values affect the values of the sought parameters.

1. Feature Scaling
2. Learning Rate
3. Error and # of iterations
4. **Regularization (L2)**



$$L = \frac{1}{2m} \left[ \sum_{i=1}^m (h^{(i)} - y^{(i)})^2 + \lambda \sum_{j=1}^n \phi_j^2 \right] \Rightarrow \min.$$

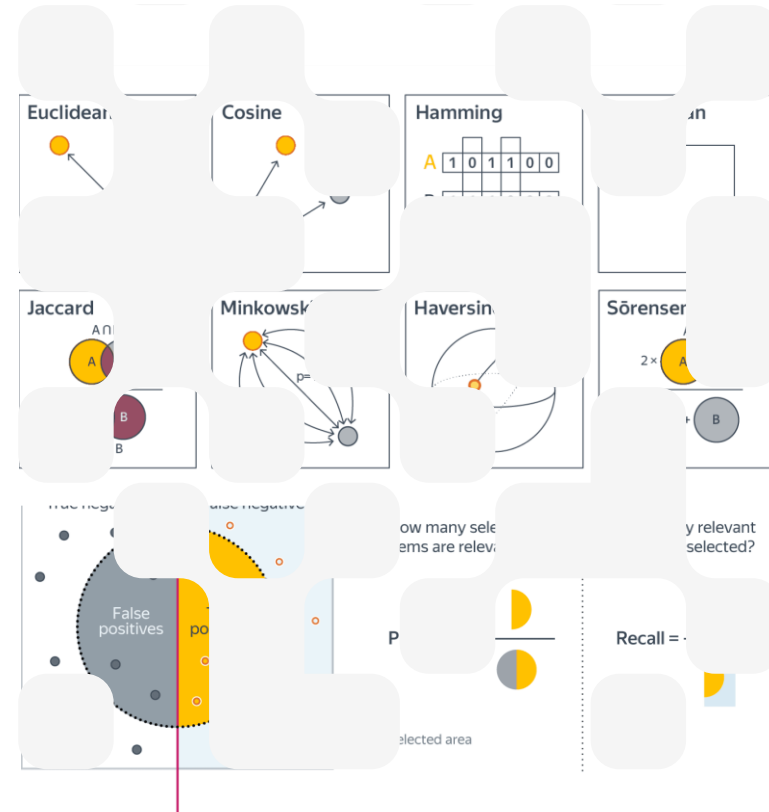
$$h(x) = \theta_j x^j, (j = 0, \dots, d)$$



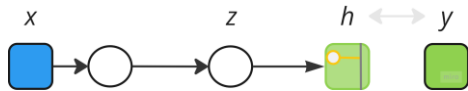
# Agenda

- I. Quality metrics in ML
- II. Data splitting, cross-validation
- III. Regularization
- IV. Batches
- V. Metric approach in ML: k-nearest neighbors (KNN)

All models are wrong, but some are useful.  
/George Box/

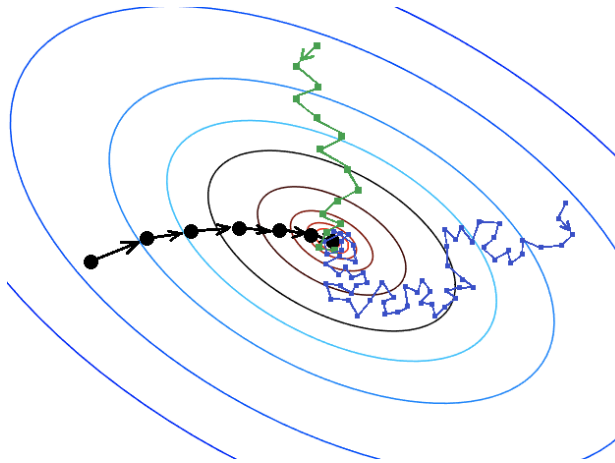


# Overfitting intuition and data splitting



Model predicts output  $h$  given input  $x$

Model *parameters* are determined during the solution of the ML problem, e.g. model weights. *Hyperparameters* are set by the user, usually not in a single way, and their values affect the values of the sought parameters.



## Batch GD

- Slowest
- Perfect gradient

## Stochastic GD

- Fastest
- Rough-estimate grad

## Mini-batch GD

- Compromise

<https://dragonnotes.org/MachineLearning/Optimization>

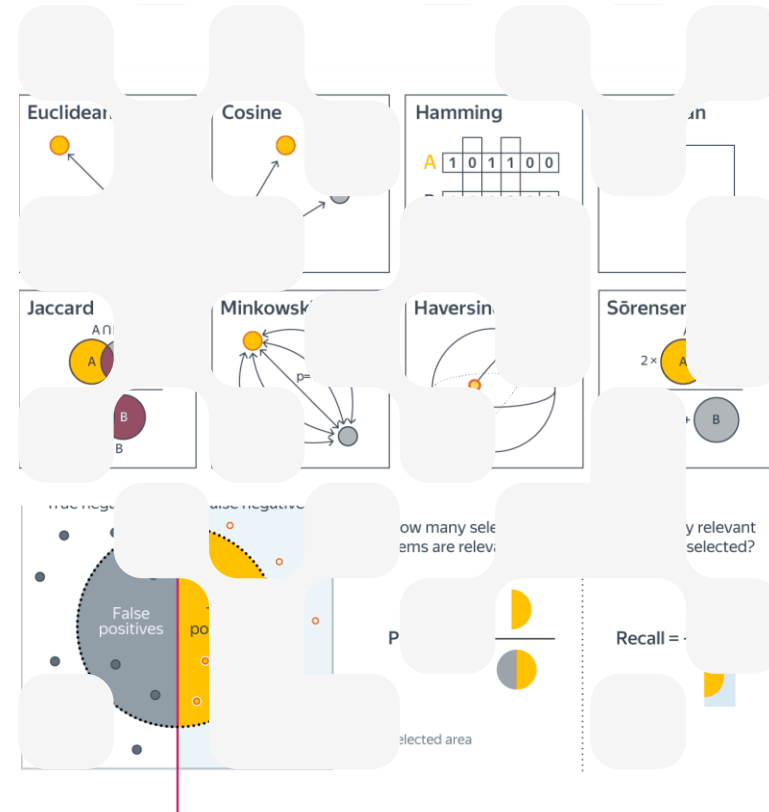


Benedict Minibatch

# Agenda

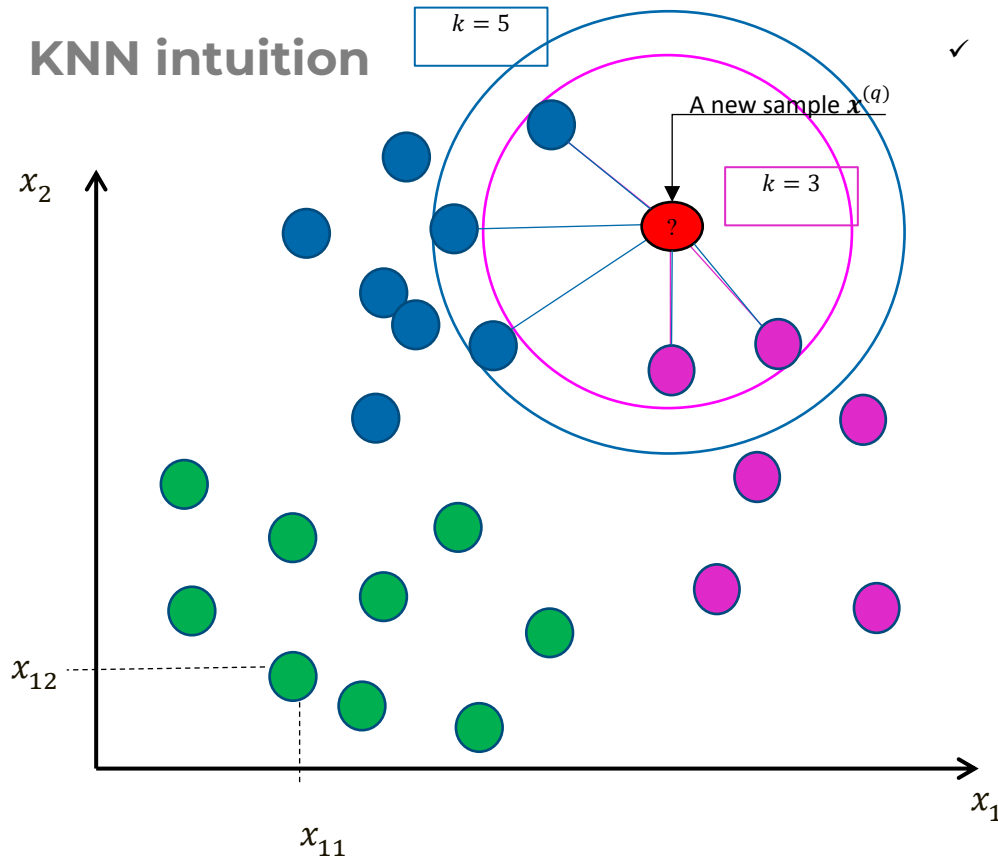
- I. Quality metrics in ML
- II. Data splitting, cross-validation
- III. Regularization
- IV. Batches
- V. Metric approach in ML: k-nearest neighbors

All models are wrong, but some are useful.  
/George Box/





# KNN intuition



- ✓ The compactness hypothesis: the assumption that similar objects are much more likely to be in the same class than in different ones.

## Training:

- The training dataset is preserved  $\{\mathbf{x}^{(i)}, y^{(i)}\}$ ;

## Classification of a new object (sample):

- Compute the distance from all samples to the new sample  $\mathbf{x}^{(q)}$ ;
- Sort the objects in ascending order of their distance to  $\mathbf{x}^{(q)}$ :

$$\rho(\mathbf{x}^{(i)}, \mathbf{x}^{(q)}) \leq \dots \leq \rho(\mathbf{x}^{(j)}, \mathbf{x}^{(q)}) \leq \dots \leq \rho(\mathbf{x}^{(r)}, \mathbf{x}^{(q)})$$

- Select the first k samples (k nearest neighbors):

$$\{\mathbf{x}^{(i)}, \dots, \mathbf{x}^{(j)}\}$$

- Assign the new sample the model class (most frequent class) among the KNN:

$$h(\mathbf{x}^{(q)}) = \operatorname{argmax}_{y_{cl} \in \mathbf{y}} \sum_{i=1}^k [y_i == y_{cl}].$$

Thank you for your attention!

[a.kornaev@innopolis.ru](mailto:a.kornaev@innopolis.ru), [@avkornaev](#)



